

Concurrent Skip Lists

Companion slides for
The Art of Multiprocessor
Programming
by Maurice Herlihy & Nir Shavit

Set Object Interface

- Collection of elements
- No duplicates
- Methods
 - add() a new element
 - remove() an element
 - contains() if element is present



Many are Cold but Few are Frozen

- Typically high % of contains() calls
- Many fewer add() calls
- And even fewer remove() calls
 - 90% contains()
 - 9% add()
 - 1% remove()
- Folklore?
 - Yes but probably mostly true




Concurrent Sets

- Balanced Trees?
 - Red-Black trees, AVL trees, ...
- Problem: no one does this well ...
- ... because rebalancing after `add()` or `remove()` is a global operation




Skip Lists

- Probabilistic Data Structure
- No global rebalancing
- Logarithmic-time search




Skip List Property

- Each layer is sublist of lower-levels




Skip List Property

- Each layer is sublist of lower-levels




Skip List Property

- Each layer is sublist of lower-levels




Skip List Property

- Each layer is sublist of lower-levels




Skip List Property

- Each layer is sublist of lower-levels
- Lowest level is entire list





Skip List Property



- Each layer is sublist of lower-levels
- Not easy to preserve in concurrent implementations ...



Search





Search





BROWN



Search




Search




Search



Search




Logarithmic



Why Logarithmic

- Property: Each pointer at layer i jumps over roughly 2^i nodes
- Pick node heights randomly so property guaranteed probabilistically



Find() -- Sequential

```
int find(T x, Node<T>[] preds, Node<T>[] succs) {  
    ...  
}
```



Find() -- Sequential

```
int find(T x, Node<T>[] preds, Node<T>[] succs) {
```

object height
(-1 if not there)



Find() -- Sequential

```
int find(T x, Node<T>[] preds, Node<T>[] succs) {
```

object height
(-1 if not there)

Object sought



Find() -- Sequential

```
int find(T x, Node<T>[] preds, Node<T>[] succs) {
```

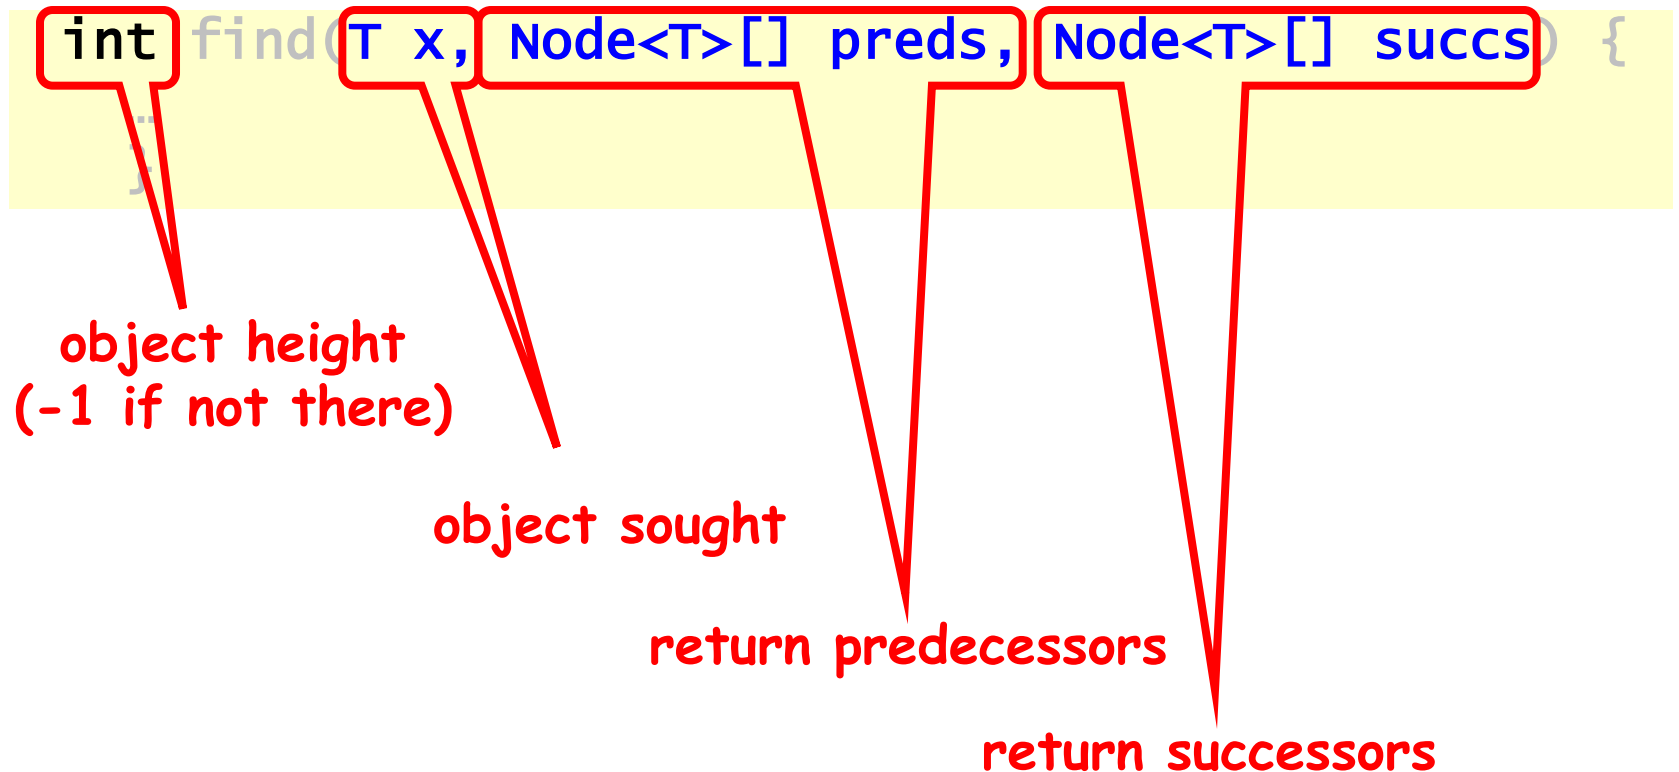
Object height
(-1 if not there)

object sought


return predecessors



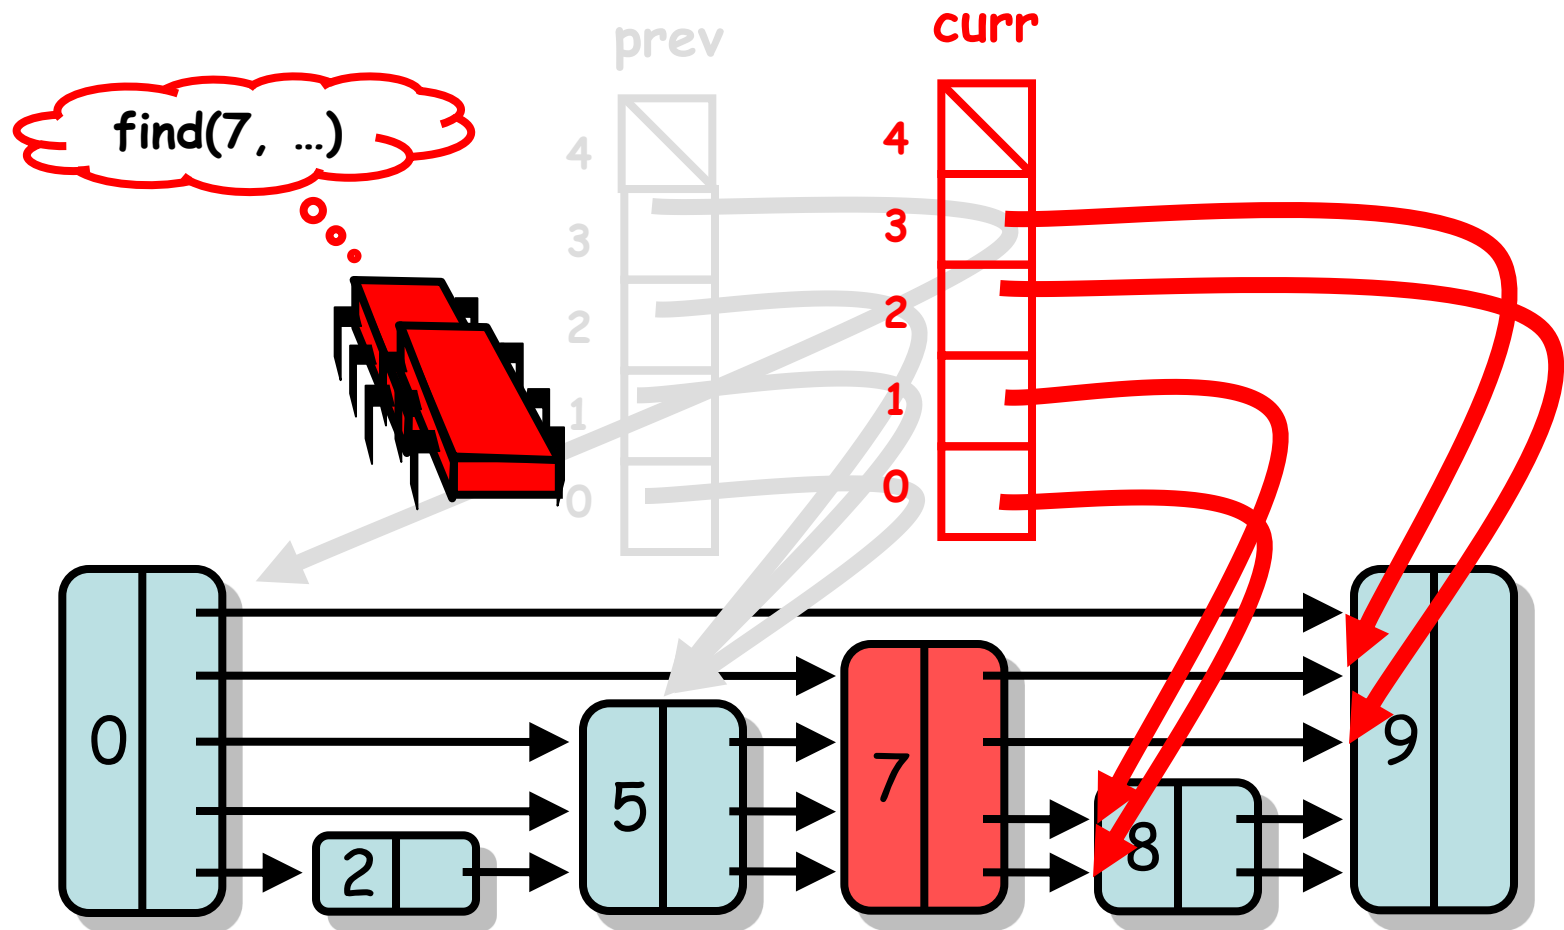
Find() -- Sequential




Successful Search




Successful Search



Unsuccessful Search



Unsuccessful Search




Lazy Skip List

- Mix blocking and non-blocking techniques:
 - Use optimistic-lazy locking for `add()` and `remove()`
 - Wait-free `contains()`
- Remember: typically lots of `contains()` calls but few `add()` and `remove()`

Review: Lazy List Remove



Review: Lazy List Remove




Present in list




BROWN

Review: Lazy List Remove



Review: Lazy List Remove




Physically deleted



BROWN



Lazy Skip Lists

- Use a mark bit for logical deletion





add(6)

- Create node of (random) height 4



add(6)



- find() predecessors



BROWN

add(6)



- find() predecessors
- Lock them



BROWN



add(6)

- **find()** predecessors
 - Lock them
 - Validate
- } **Optimistic approach**



add(6)


- find() predecessors
- Lock them
- Validate
- Splice




BROWN

add(6)

- find() predecessors
- Lock them
- Validate
- Splice
- Unlock




remove(6)




remove(6)

- find() predecessors



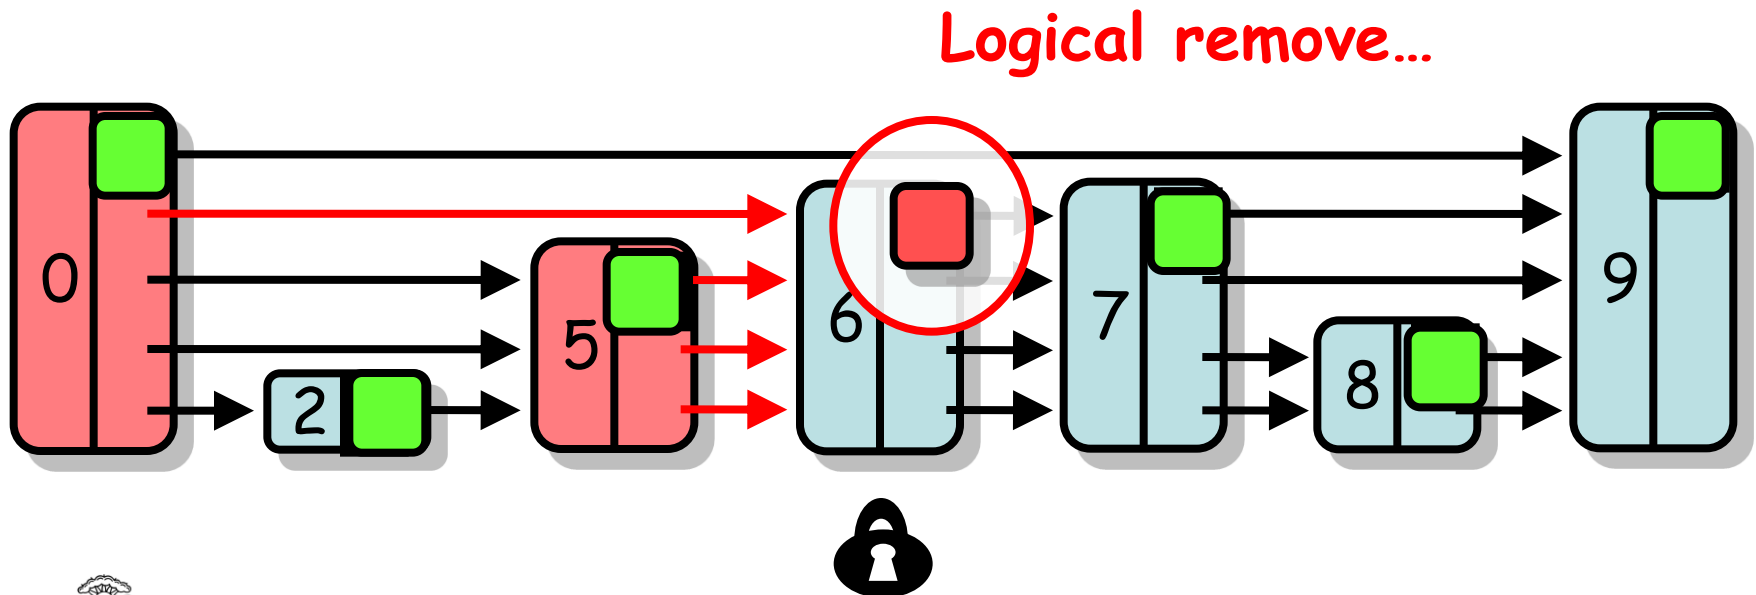
remove(6)

- find() predecessors
- Lock victim




remove(6)

- find() predecessors
- Lock victim
- Set mark (if not already set)




remove(6)

- find() predecessors
- Lock victim
- Set mark (if not already set)
- Lock predecessors (ascending order) & validate




remove(6)

- find() predecessors
- Lock victim
- Set mark (if not already set)
- Lock predecessors (ascending order) & validate
- Physically remove




remove(6)

- find() predecessors
- Lock victim
- Set mark (if not already set)
- Lock predecessors (ascending order) & validate
- Physically remove




remove(6)

- find() predecessors
- Lock victim
- Set mark (if not already set)
- Lock predecessors (ascending order) & validate
- Physically remove




remove(6)

- find() predecessors
- Lock victim
- Set mark (if not already set)
- Lock predecessors (ascending order) & validate
- Physically remove




remove(6)

- find() predecessors
- Lock victim
- Set mark (if not already set)
- Lock predecessors (ascending order) & validate
- Physically remove




contains(8)

- Find() & not marked




contains(8)

Node 6 removed while traversed




BROWN

contains(8)



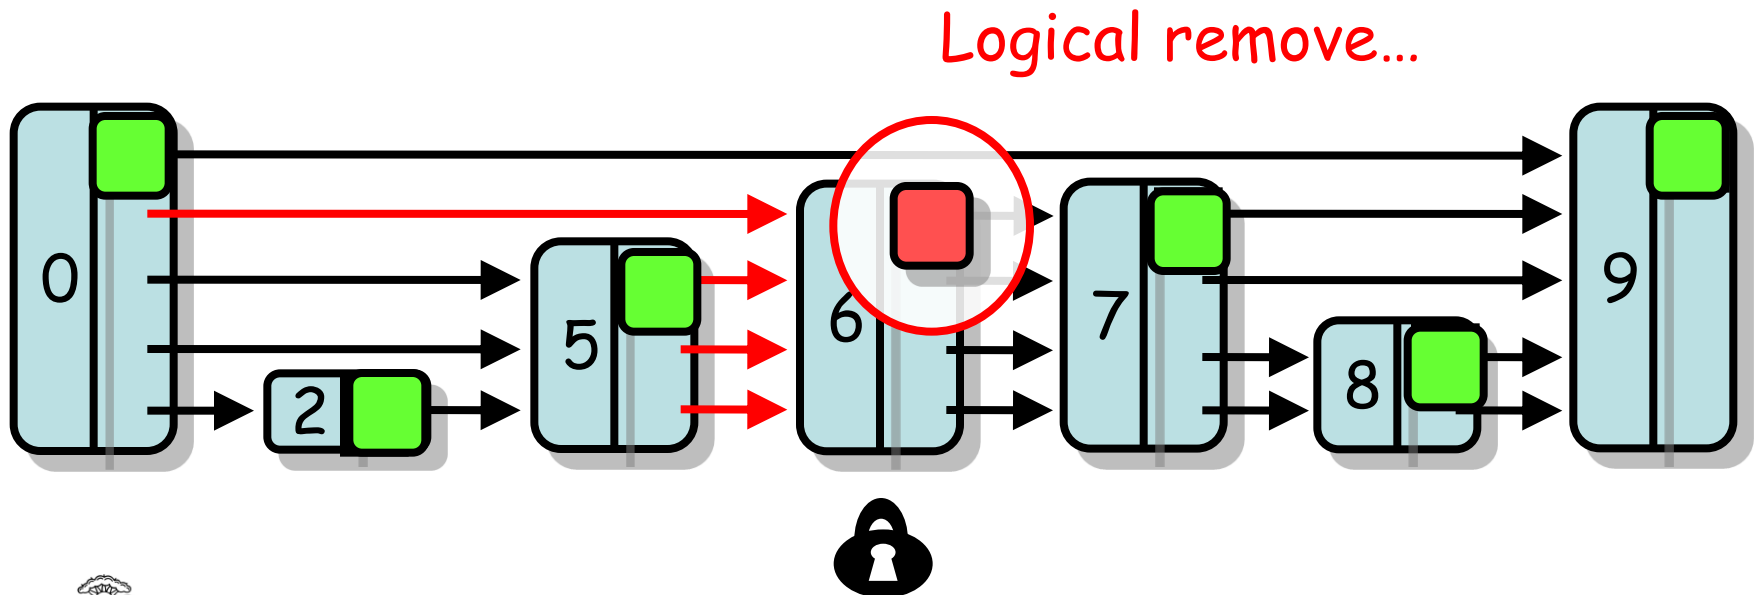
contains(8)

Prove: an unmarked
node (like 8)
remains reachable





remove(6): Linearization

- Successful remove happens when bit is set



Add: Linearization


- Successful add() at point when fully linked
- Add fullyLinked bit to indicate this
- Bit tested by contains()



Unsuccessful add(6)


Linearization

- When not fully-linked unmarked node found
- Pause while fullyLinked bit sets



contains(7): Linearization


- When do we linearize unsuccessful Search?



contains(7): Linearization

- When do we linearize unsuccessful Search?


But what if a new
7 added concurrently?



contains(7): Linearization

- When do we linearize unsuccessful Search?

Prove: at some point 7 was not in the skip list




Look-Free Skip List

- We can't manipulate multiple reference at the same time
 - Can't maintain the SkipList property
 - Each list could be **not** a sublist of the list at levels below
- The Abstract set is defined by the bottom-level
 - A key is in the set if there is a node with that key whose next reference is unmarked in the bottom level list




LockFree Skip List

- Use a mark bit for logical deletion
- Did I mark it ?



LockFree Skip List

- Use a mark bit for logical deletion
- Use a mark bit for each linked list



How to Add/Remove ?

- Each level of the list is treated as a LockFree list
- Use *CAS()* to insert
- Mark the next reference to remove
- *find()* method cleans up marked nodes
 - It's possible for a node to be physically removed while it's being linked at lower level



Add()

- Call find()
 - to determine whether a node is already in the list, find its set of predecessors and successors
- A new node is prepared
 - with randomly chosen topLevel, and its next references are directed to the potential successors
- Add to the bottom level LockFreeList
 - Successively links the node in higher levels


Linearization point




add(12)




find(12)



Remove(2)



Remove(9)




Remove(15)




BROWN

Art of Multiprocessor
Programming© Herlihy Shavit
2007

add(12)



Remove(2)




Remove(9)




BROWN

Art of Multiprocessor
Programming© Herlihy Shavit
2007

add(12)



Remove(2)



Remove(9)




insert(12)




BROWN

add(12)



Remove(2)



Remove(9)




insert(12)




BROWN

add(12)



Remove(2)



Remove(9)




insert(12)




BROWN

add(12)



Remove(2)




Remove(9)




BROWN

Art of Multiprocessor
Programming© Herlihy Shavit
2007

add(12)



Remove(2)





Remove(9)



BROWN

Art of Multiprocessor
Programming© Herlihy Shavit
2007

add(12)



Remove(9)

Art of Multiprocessor
Programming© Herlihy Shavit
2007



BROWN

Remove()

- Call find()
 - to determine whether an unmarked node with the target key is in the **bottom** level list
- Logical remove
 - mark from the top level to the bottom level
- physical remove
 - is done either by itself or by the other find() traversing the link



contains()

- Can't use find()
 - trying to remove marked nodes might generate too much contention
- Can't use LockFreeList's contains()
 - might skip nodes reachable from the bottom level
- Use find()
 - without actually doing physical remove




contains(18)




contains(18)

→ current
→ prev



Remove(9)




Remove(15)




BROWN

contains(18)




Contains(18)

→ current
→ prev



Remove(9)



Remove(15)




BROWN

contains(18)




contains(18)

→ current
→ prev



Remove(9)




Remove(15)




BROWN

contains(18)

→ current
→ prev



Remove(9)



Remove(15)



BROWN

contains(18)




Contain(18




current



prev



Remove(9)



Remove(15)



BROWN


Art of Multiprocessor
Programming© Herlihy Shavit
2007

contains(18)




contains(18)

→ current
→ prev



Remove(9)



Remove(15)

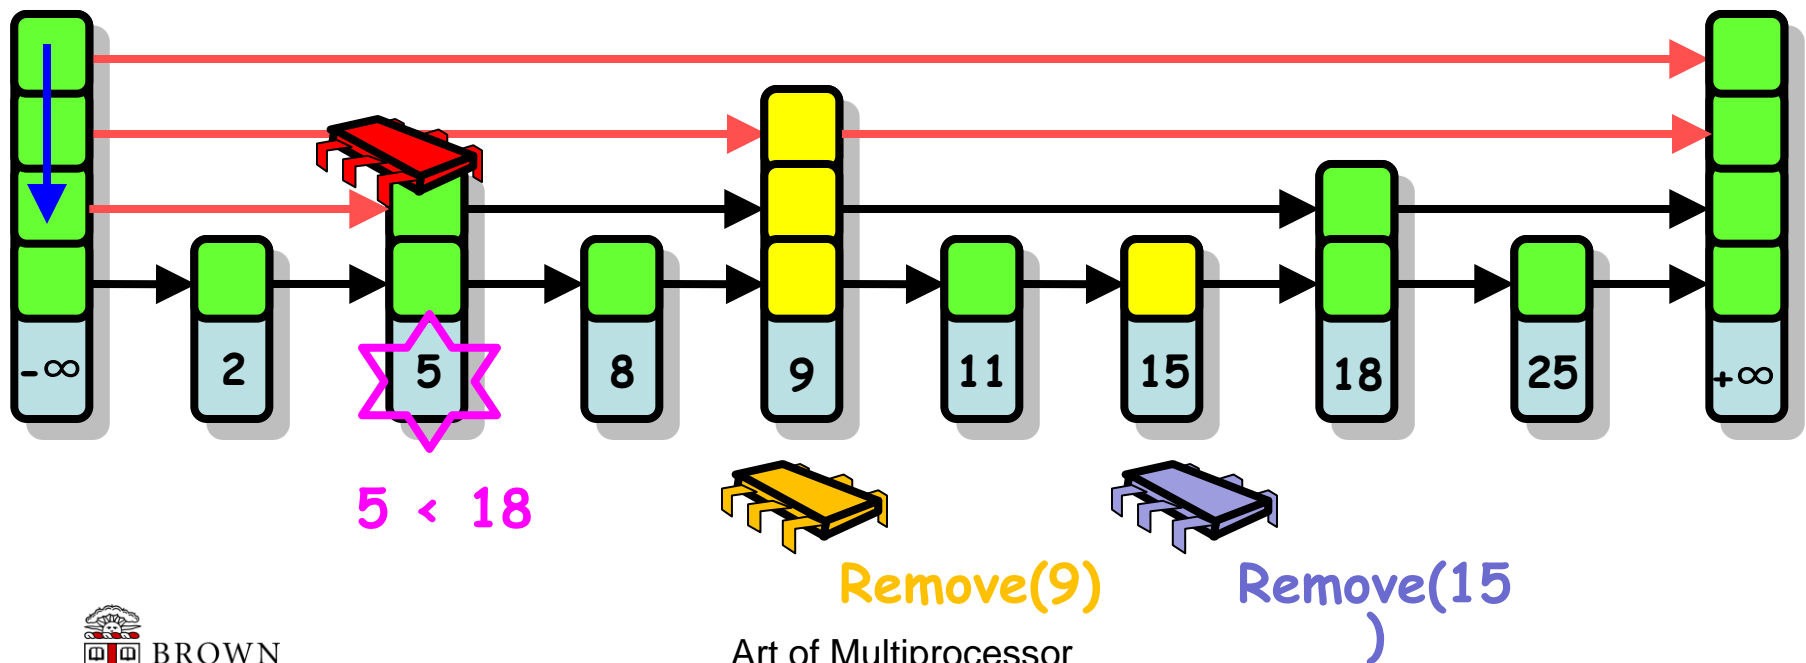


BROWN

contains(18)

Contains(18)
)

→ current
→ prev




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)




BROWN

contains(18)

**Contain(18
)**

→ current
→ prev



BROWN


Art of Multiprocessor
Programming© Herlihy Shavit
2007

contains(18)




contains(18)

→ current
→ prev



Remove(9)



Remove(15)




BROWN

contains(18)

Contains(18)

→ current
→ prev




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)




BROWN

contains(18)

Contain(18) returns TRUE

→ current
→ prev



BROWN

When is contains() linearized ?


- It might be possible contains() returns false even though the item has been inserted !

contains(18)




Contain(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN

contains(18)




Contain(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN

contains(18)




Contain(18)

→ current
→ prev



Remove(9)



Remove(15)




Remove(18)




BROWN

contains(18)

→ current
→ prev



Remove(9)



Remove(15)




Remove(18)




BROWN

contains(18)




contains(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)

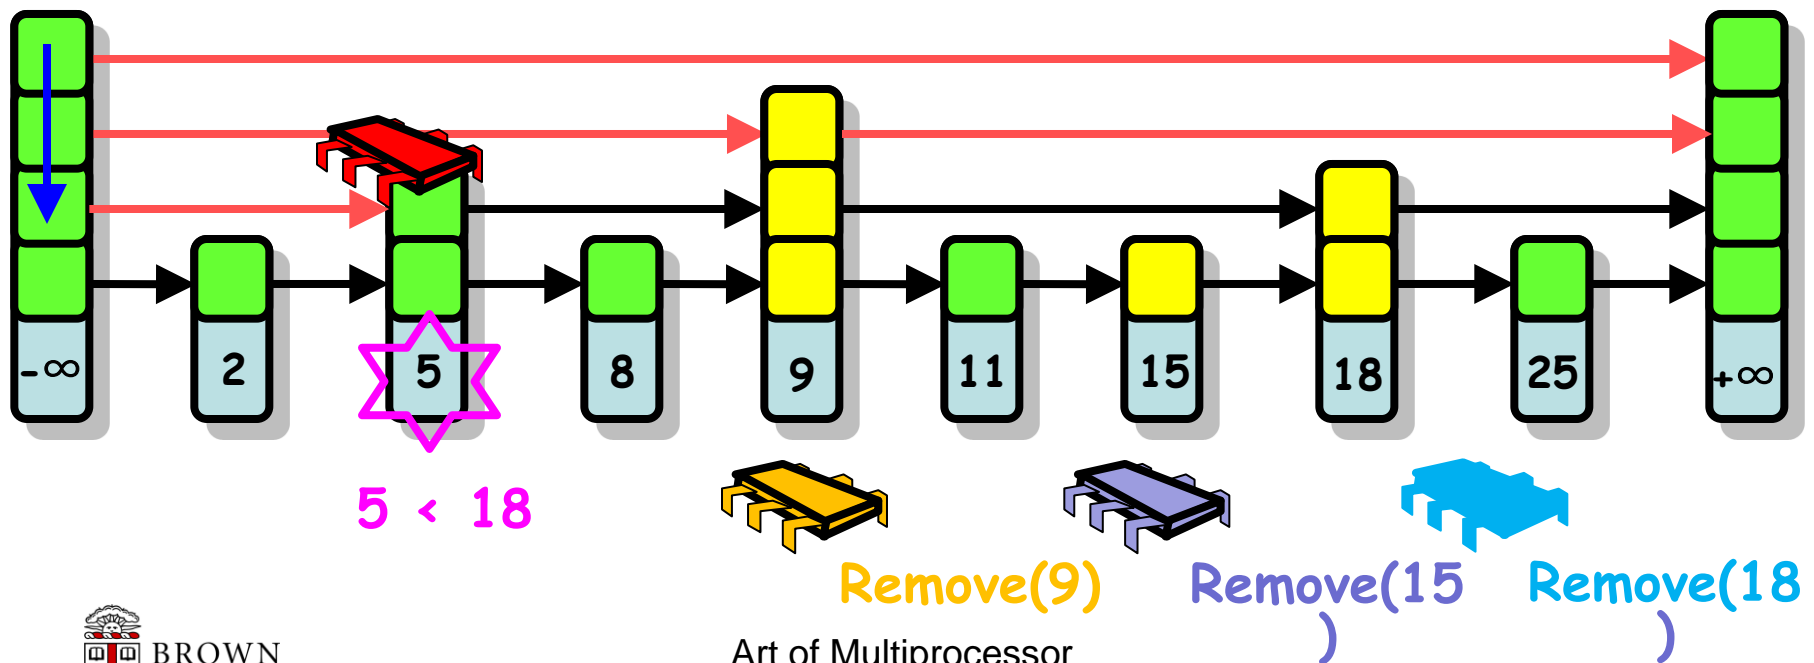


BROWN

contains(18)

Contain(18
)

→ current
→ prev




BROWN

contains(18)

Contains(18)

→ current
→ prev




BROWN

contains(18)

**Contain(18
)**

→ current

→ prev



Remove(9)

Remove(15)



Remove(18)




BROWN

Art of Multiprocessor
Programming© Herlihy Shavit
2007


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)




Remove(18)




BROWN

contains(18)




contains(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)

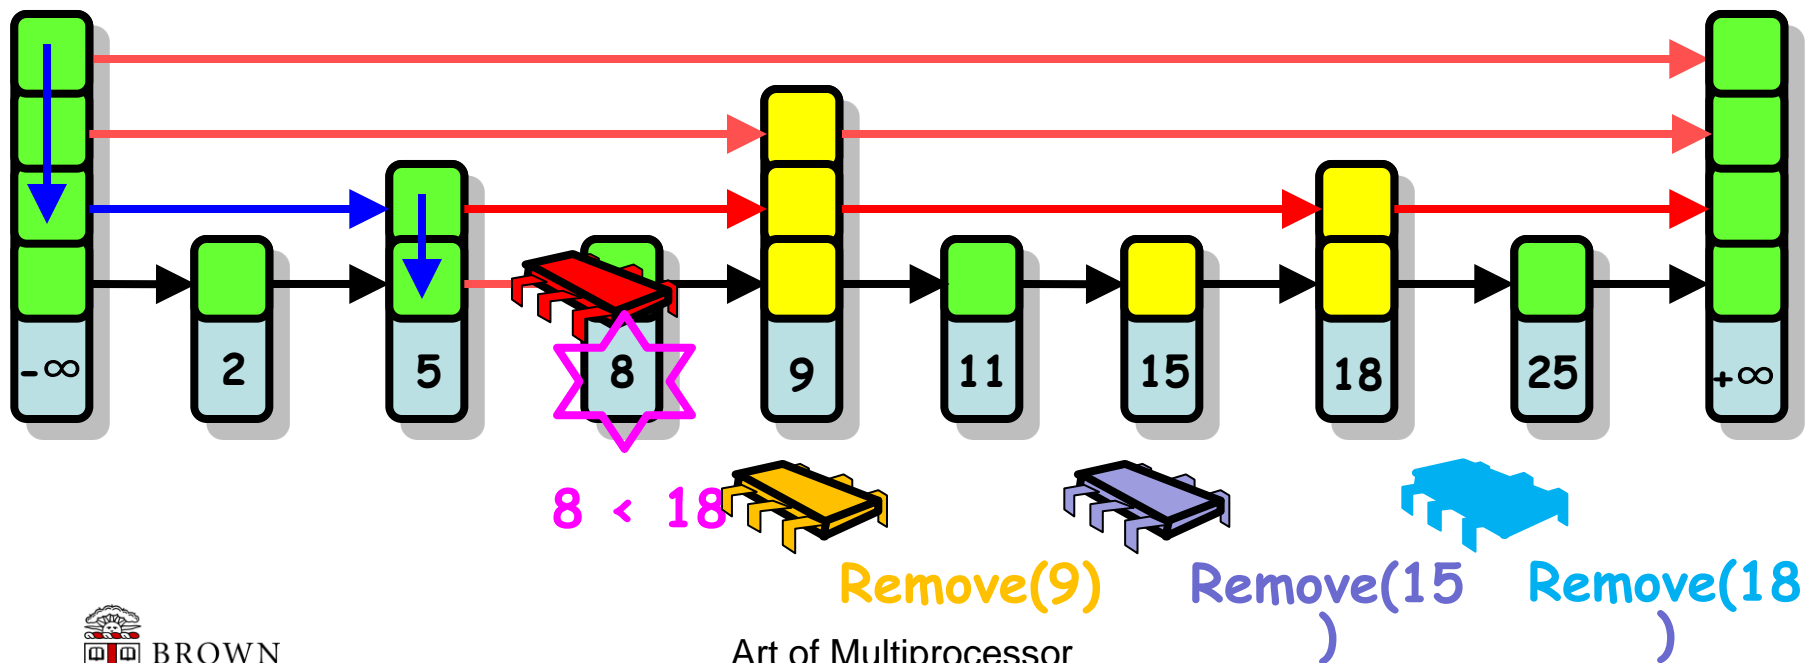


BROWN

contains(18)

Contains(18)

→ current
→ prev




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN


contains(18)

Contains(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN


contains(18)

Contain(18
)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN


contains(18)

Contain(18)
)
Add(18)
)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN

contains(18)


Contain(18
)

Add(18) -> find(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN

contains(18)


Contain(18
)

Add(18) -> find(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN

contains(18)


Contain(18
)

Add(18) -> find(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN

contains(18)


Contain(18
)

Add(18) -> find(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN

contains(18)


Contain(18
)

Add(18) -> find(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN

contains(18)


Contain(18
)

Add(18) -> find(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)




BROWN

contains(18)


Contain(18)
)

Add(18) -> find(18)

→ current
→ prev



Remove(9)



Remove(15)



Remove(18)



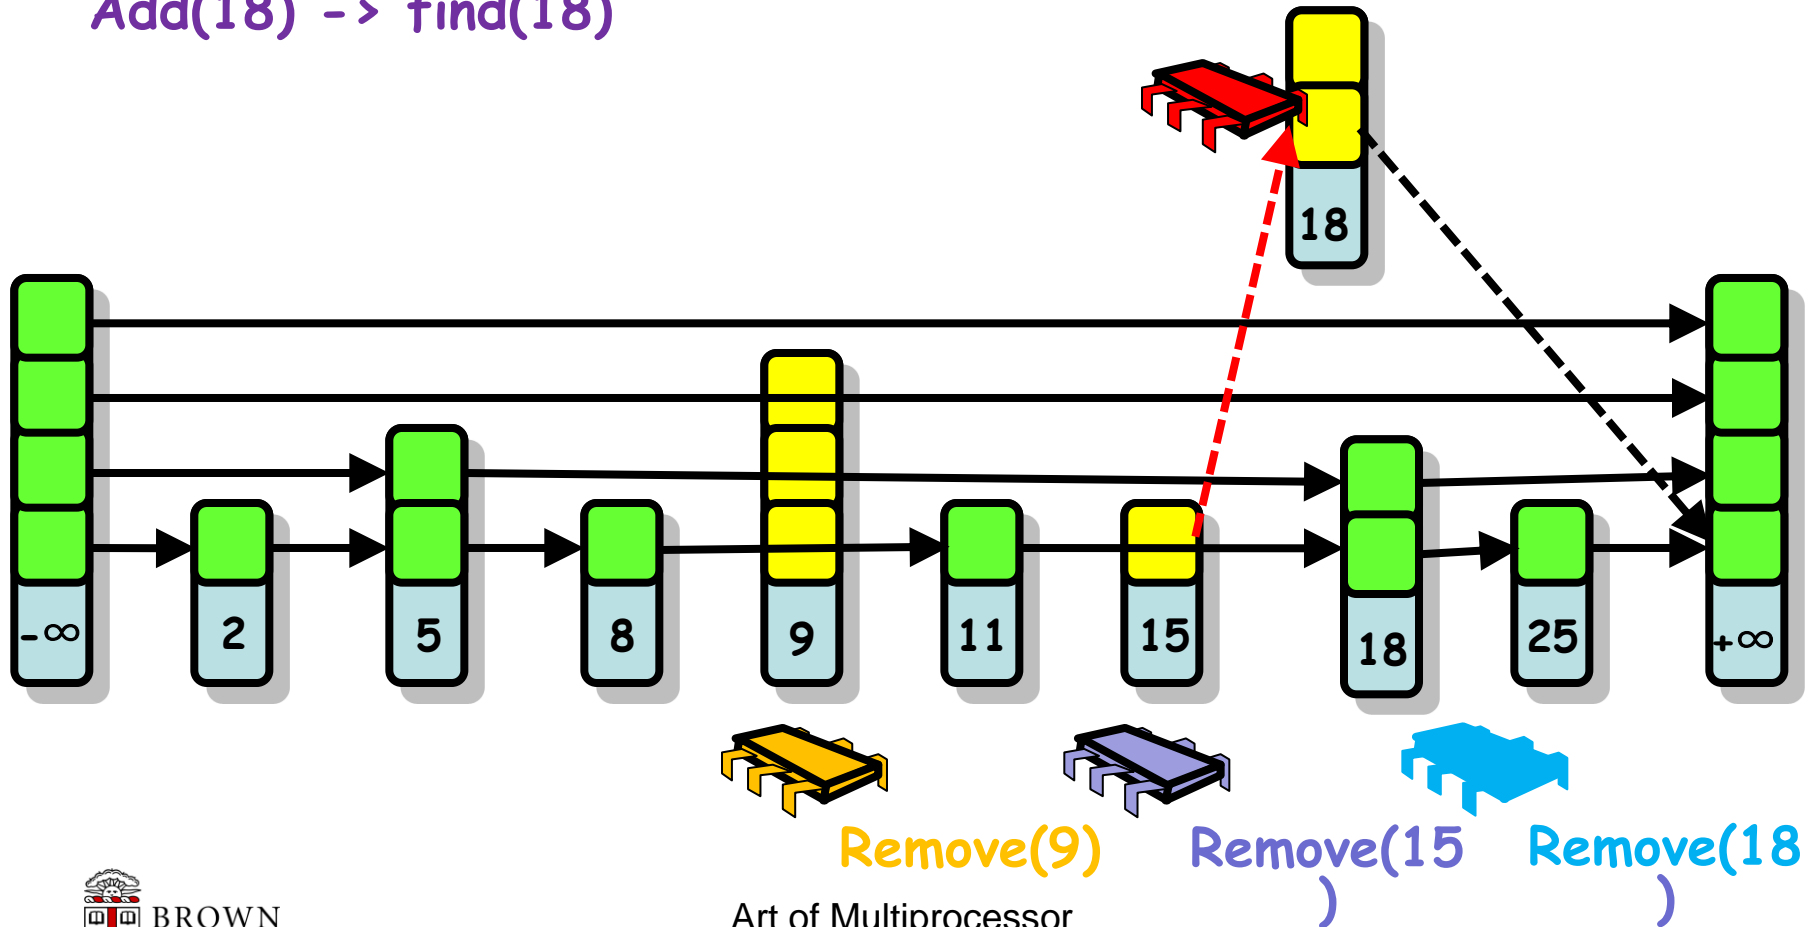
BROWN

contains(18)

Contain(18
)

Add(18) -> find(18)

→ current
→ prev



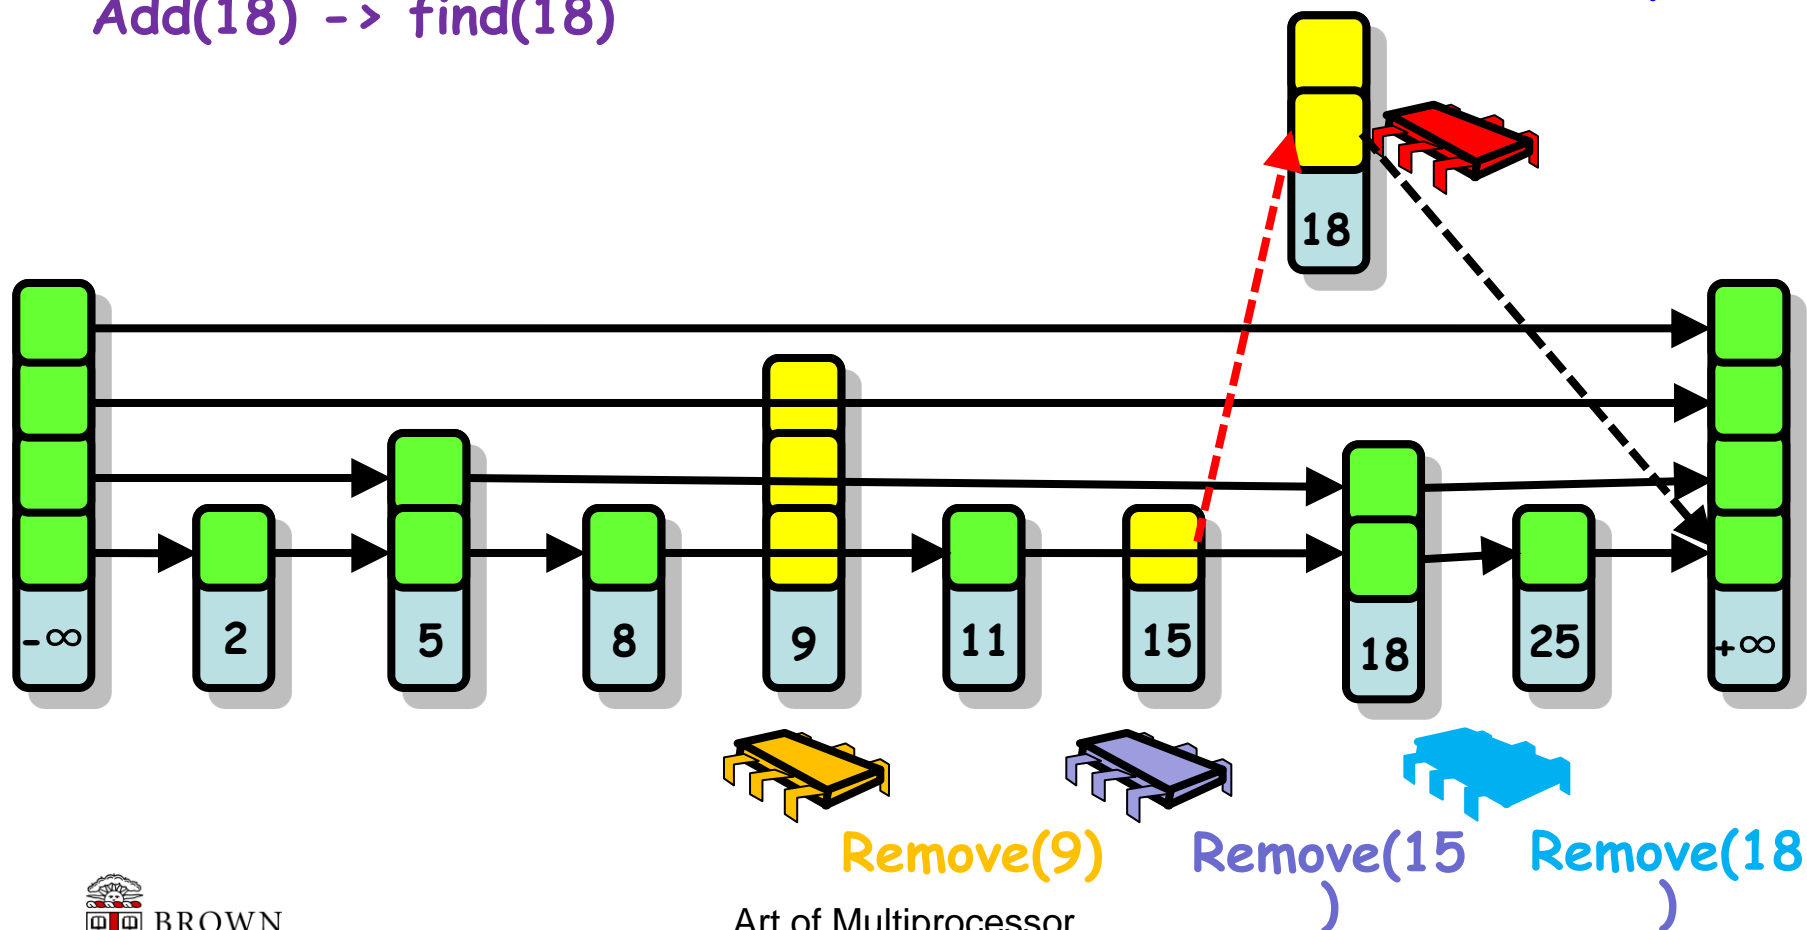
BROWN

contains(18)

Contain(18) returns FALSE !

Add(18) -> find(18)

→ current
→ prev




BROWN


A Simple Experiment

- Each thread runs 1 million iterations, each either:
 - `add()`
 - `remove()`
 - `contains()`
- Item and method chosen in random from some distribution


Lazy Skip List: Performance



Lazy Skip List: Performance



Lazy Skip List: Performance



Summary

- Lazy Skip List
 - Optimistic fine-grained Locking
- Performs as well as the lock-free solution in "common" cases
- Simple

This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](https://creativecommons.org/licenses/by-sa/3.0/).

- **You are free:**
 - **to Share** — to copy, distribute and transmit the work
 - **to Remix** — to adapt the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
 - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
 - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.

